

What you missed.

- Keynotes: Not really interesting, none were very good at presenting
- Full Papers: A couple of really good papers and presentations, but more lesser quality
- Short Papers/Posters: Not much “IR” going on here; ML, NLP, industry focused.
- Tutorials: The one I went to was interesting;
Workshop: Skippable
- Food: People got sick, but banquet was good

scells.me/cikm2017

My top 5 papers.

1. A Neural Candidate-Selector for Automatic Structured Clinical Text Annotation



2. Active Sampling for Large-scale Information Retrieval Evaluation



3. DeepRank: A New Deep Architecture for Relevance Ranking in Information Retrieval

4. Adaptive Persistence for Search Effectiveness Measures

5. FA*IR: A Fair Top-k Ranking Algorithm

scells.me/cikm2017