

Learning to Rank Citations

- What's been done?
 - Unlike active learning, which requires **explicit relevance feedback** in order to **re-rank** a result set, learning to rank seeks to **learn a model** capable of ranking documents [84]
 - It does **not rely** on explicit relevance feedback
 - Existing learning to rank algorithms have been shown to **outperform** previous probabilistic ranking functions [15]
- What's the gap?



GAP

Unclear how to learn a model for this domain.

[84] Andrew Trotman. Learning to rank. Information Retrieval, 8(3):359–381, 2005.

[15] Olivier Chapelle and Yi Chang. Yahoo! learning to rank challenge overview. In Proceedings of the Learning to Rank Challenge, pages 1–24, 2011.

Screening Prioritisation

- Any active/machine learning approach to screening prioritisation still requires **screening all citations**
- In some cases it is only used to **replace second screener**
- Automatically ranking citations
 - **No explicit relevance feedback**
 - **Cutoff estimation**